

# WealthPilot

## Business Case — AI Co-Pilot for Independent Financial Advisors

### The User

The target customer is the independent Registered Investment Advisor (RIA) managing between \$50M and \$500M in assets under management — specifically the solo practitioner or small team (1–5 people) who left a wirehouse or large institution to run their own book.

This advisor manages 60–120 clients with no quant team, no risk desk, and no Bloomberg terminal. Their toolkit is Excel, a custodian portal (Schwab, Fidelity, or Pershing), and a CRM. They are experts at client relationships — not at running portfolio analytics at 7am every morning.

**Concrete persona:** Jane Smith, CFA. 20 years at Merrill Lynch, went independent 4 years ago. Manages 82 clients, \$180M AUM. Spends Sunday evenings manually reviewing drift reports. Misses an average of \$8,400/year in tax loss harvesting savings per client due to infrequent review.

### The Problem

Jane's Monday morning routine: export 82 account CSVs from Schwab, paste into Excel, run VLOOKUP to check drift vs. targets, manually verify each proposed trade against the client's Investment Policy Statement (IPS), calculate whether harvestable losses exceed trading costs, and draft client emails explaining any changes — one at a time.

This process consumes 12–15 hours per month before she touches actual client work. Because it is so expensive in time, the review that should happen nightly only happens quarterly. The result: missed rebalancing windows, foregone tax alpha, and compliance exposure.

The institutional tools that solve this problem — Bloomberg, Orion, Black Diamond — cost **\$150,000–\$250,000 per year**. WealthPilot delivers the same morning intelligence workflow at **\$299/month**.

### The Economics

Pricing is \$299/month per advisor (\$2,990/year on annual plan). The unit economics are compelling:

| Cost Item                                      | Monthly Cost |
|------------------------------------------------|--------------|
| Gemini 2.0 Flash tokens (80 clients × 30 days) | \$3.20       |
| GCP Cloud Run compute                          | \$6.00       |
| Storage, logging, monitoring                   | \$1.50       |
| Support allocation                             | \$2.30       |
| Total COGS                                     | \$13.00      |
| Revenue (price per advisor)                    | \$299.00     |
| Gross Margin                                   | 95.7%        |

**Where it breaks:** If token costs increase 10× (unlikely given current Gemini pricing trends), gross margin drops to ~83% — still healthy. The model is unviable below \$50 revenue per advisor; our price point has 6× headroom. Primary churn risk is advisors who trial but don't integrate it into their daily workflow — mitigated by onboarding support in Phase 1.

**TAM:** ~15,000 independent RIAs in the US (SEC registration data). Addressable initially: ~4,000 (AUM \$50M–\$500M, tech-forward). At 1% penetration → 40 advisors → **\$143K ARR**. At 5% → \$718K ARR (seed-stage viable).

Why These Technical Choices?

| Concept                   | Application in WealthPilot                                                |
|---------------------------|---------------------------------------------------------------------------|
| Multi-Agent Orchestration | Three specialized agents coordinated by orchestrator.py                   |
| Structured Output         | Every agent returns typed, validated dicts — machine-readable audit trail |
| LLM Tool Use / Prompting  | Gemini called with per-client context for personalized email drafting     |
| Parallel Execution        | Portfolio monitor processes all client accounts in single pipeline pass   |
| RAG                       | Client IPS documents retrieved per-client and injected into LLM prompt    |
| Human-in-the-Loop         | Advisor must approve all rebalance orders and emails before execution     |

**Multi-Agent Architecture:** The problem decomposes naturally into three specialist tasks with different data needs and risk tolerances. Portfolio drift calculation is deterministic math; suitability checking is rule-based compliance; email drafting is creative personalization. A monolithic LLM call would conflate them. Separating into agents means each can be tested, audited, and improved independently — critical for a regulated financial product.

**Structured Output at Every Stage:** Every agent produces typed, validated dictionaries — not free-form text. Regulators examining a suitability decision need a machine-readable audit trail. Structured output also makes the human-in-the-loop approval step safe: the advisor approves a specific trade object, not a vague AI recommendation.

**RAG for Suitability:** Each client's IPS is a firm-specific document. We cannot bake 80 client IPS files into a system prompt. The suitability agent retrieves the relevant client's IPS and injects it into context at query time — retrieval-augmented generation applied directly to compliance.

**Gemini via Vertex AI:** Cost per token is the dominant unit economics variable. Gemini 2.0 Flash is among the most cost-efficient capable models available. Vertex AI provides enterprise security and data residency guarantees that a financial services product requires — client portfolio data cannot transit a consumer API.

**Human-in-the-Loop Design:** No trade executes and no email sends without advisor approval. This is not only a UX choice — it is a regulatory requirement. RIAs carry fiduciary duty. WealthPilot positions itself as an AI co-pilot, not an autonomous trader. This reduces regulatory risk, builds advisor trust, and differentiates clearly from robo-advisors competing in a different segment.

Go-To-Market

- Phase 1 (0–6 months): Direct outreach to NAPFA members. Free 90-day trial for first 20 advisors. Gather workflow feedback, tighten product-market fit.
- Phase 2 (6–18 months): Custodian partnerships via Schwab and Fidelity RIA channels. Integration into their advisor portal as a recommended intelligence tool.
- Phase 3 (18–36 months): API layer for portfolio management platforms (Orion, Nitrogen/Riskalyze). WealthPilot becomes the intelligence layer that plugs into existing workflows.

"We're not replacing the advisor. We're giving them their Sunday evenings back."